

POLIST_02_notebookessais

September 8, 2021

```
[1]: import numpy as np
import pandas as pd

from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_samples, silhouette_score, \
    adjusted_rand_score

import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import matplotlib.cm as cm

%matplotlib inline
import seaborn as sns

import datetime
from dateutil.relativedelta import *
import scipy.stats as stats
import math as mt
import timeit

from sklearn import decomposition, preprocessing

import plotly.graph_objects as go
import plotly.express as px
```

```
Reading input data customers = pd.read_csv("./data/olist_customers_dataset.csv")
geoloc = pd.read_csv("./data/olist_geolocation_dataset.csv") items
= pd.read_csv("./data/olist_order_items_dataset.csv") payments
= pd.read_csv("./data/olist_order_payments_dataset.csv") reviews
= pd.read_csv("./data/olist_order_reviews_dataset.csv") orders =
pd.read_csv("./data/olist_orders_dataset.csv")
```

```
[2]: customers = pd.read_csv("./data/customers.csv")
#geoloc = pd.read_csv("./data/olist_geolocation_dataset.csv")
items = pd.read_csv("./data/items.csv")
payments = pd.read_csv("./data/payments.csv")
reviews = pd.read_csv("./data/reviews.csv")
```

```
orders = pd.read_csv("../data/orders.csv")
```

```
[3]: #convert ['order_purchase_timestamp', 'order_approved_at',  
→ 'order_delivered_carrier_date', 'order_delivered_customer_date',  
→ 'order_estimated_delivery_date'] to datetime  
date_columns = ['order_purchase_timestamp', 'order_approved_at',  
→ 'order_delivered_carrier_date', 'order_delivered_customer_date',  
→ 'order_estimated_delivery_date']  
for col in date_columns:  
    orders[col] = pd.to_datetime(orders[col])
```

```
[24]: # function transforming timedelta values to seconds  
def to_seconds(x):  
    return x.total_seconds()
```

0.0.1 Constitution du score RFM

Data aggregation

```
[4]: #Processing related to "orders"  
  
order_custom = orders.merge(customers[['customer_id', 'customer_unique_id']],  
→ on='customer_id')  
part1 = order_custom[['order_purchase_timestamp', 'customer_unique_id']].  
→ groupby(['customer_unique_id']).max()  
# test_2c = part1  
  
#Processing related to "customers"  
  
part2 = customers[['customer_id', 'customer_unique_id']].  
→ groupby(['customer_unique_id']).count()  
  
#Processing related to "payments"  
  
part30 = payments[['payment_value', 'order_id']].groupby(['order_id']).sum()  
order_payment = orders.join(part30, on=['order_id'])  
part31 = order_payment[['payment_value', 'customer_id']].  
→ merge(customers[['customer_unique_id', 'customer_id']], on='customer_id')  
  
part3 = part31[['payment_value', 'customer_unique_id']].  
→ groupby(['customer_unique_id']).sum()
```

```
[5]:
```

```
rfm = (part1.merge(part2, on='customer_unique_id')).merge(part3, on =
↳ 'customer_unique_id')

rfm = rfm.rename(columns={'customer_id': 'frequency', 'order_purchase_timestamp':
↳ 'recency', 'payment_value': 'monetary'})

rfm['recency'] = rfm['recency'] - rfm['recency'].max()

rfm
```

```
[5]:
```

customer_unique_id	recency	frequency	monetary
0000366f3b9a7992bf8c76cfd3221e2	-112 days +19:55:50	1	141.90
0000b849f77a49e4a4ce2b2a4ca5be3f	-115 days +20:10:50	1	27.19
0000f46a3911fa3c0805444483337064	-537 days +06:04:26	1	86.22
0000f6ccb0745a6a4b88665a16c9f078	-321 days +05:29:04	1	43.62
0004aac84e0df4da2b147fca70cf8255	-288 days +04:45:05	1	196.89
...
ffffcf5a5ff07b0908bd4e2dbc735a684	-447 days +05:59:59	1	2067.42
fffea47cd6d3cc0a88bd621562a9d061	-262 days +05:07:19	1	84.58
ffff371b4d645b6ecea244b27531430a	-568 days +00:48:39	1	112.46
ffff5962728ec6157033ef9805bacc48	-119 days +00:17:04	1	133.69
ffffd2657e2aad2907e67c3e9daecbeb	-484 days +05:18:08	1	71.56

[93356 rows x 3 columns]

```
[6]: rfm['recency'].min()
```

```
[6]: Timedelta('-714 days +21:16:01')
```

Building rfm scores - percentile-based

Function definition

```
[7]: #Cette fonction fait correspondre à toute valeur d'une variable, un rang qui
↳ correspond à l'inter-percentile auquel il
#appartient dans la distribution

def my_percentile_rank_fix(v,thr):
    if v <= thr:
        res = 2
    else:
        res = 4
    return res
```

```

#Cette fonction fait correspondre à toute valeur d'une variable, un rang qui
↳correspond à l'inter-percentile auquel il
# appartient dans la distribution
def my_percentile_rank(col, v, n):

    range_ = np.linspace(1,n-1,n-1)
    lim=np.percentile(col,range_*(100/n))
    if np.sign(min(v-lim))>0:
        res = np.argmin(abs(v-lim))+2
    else:
        res = np.argmin(abs(v-lim))+1
    return res

#Cette variante de "my_percentile_rank" utilise
def my_percentile_rank_rec(col, v, n):

    range_ = np.linspace(1,n-1,n-1)
    lim=np.percentile(col,range_*(100/n))
    if np.sign(min(v-lim))>np.timedelta64(0,'ns'):# Ici la valeur 0 est de type
↳timedelta64
        res = np.argmin(abs(v-lim))+2
    else:
        res = np.argmin(abs(v-lim))+1
    return res

#Cette fonction fait correspondre à toute valeur v d'une variable contenue dans
↳col, un rang qui correspond à l'inter-percentile auquel il
#appartient dans la distribution
def my_segment(col, v, n):

    range_ = np.linspace(1,n-1,n-1)
    lim=range_*(mt.floor(col.max()/n))
    if np.sign(min(v-lim))>0:
        res = np.argmin(abs(v-lim))+2
    else:
        res = np.argmin(abs(v-lim))+1
    return res

```

```

[8]: def my_quintile_recency_rank(v):
    return my_percentile_rank_rec(rfm['recency'],v,5)

def my_quintile_frequency_rank(v):
    # return my_percentile_rank_fix(v,rfm["frequency"].unique().mean())
    return my_percentile_rank_fix(v,1)

def my_quintile_monetary_rank(v):
    return my_percentile_rank(rfm['monetary'],v,5)

```

Actual computation

```
[9]: #Recency score
start_time = timeit.default_timer()
rfm['recency_s'] = rfm['recency'].apply(lambda x: my_quintile_recency_rank(x))
elapsed = timeit.default_timer() - start_time

print(elapsed)
```

905.4680586620016

```
[10]: #Frequency score
start_time = timeit.default_timer()
rfm['frequency_s'] = rfm['frequency'].apply(lambda x:
    ↳my_quintile_frequency_rank(x))
elapsed = timeit.default_timer() - start_time

print(elapsed)
```

0.030257911996159237

```
[11]: #Monetary score
start_time = timeit.default_timer()
rfm['monetary_s'] = rfm['monetary'].apply(lambda x:
    ↳my_quintile_monetary_rank(x))
elapsed = timeit.default_timer() - start_time

print(elapsed)
```

164.6787575740018

```
[12]: rfm['recency_s'].value_counts()
```

```
[12]: 1    26086
      2    19754
      3    19359
      5    18671
      4     9486
      Name: recency_s, dtype: int64
```

```
[13]: # Computation of a synthetic rfm score
rfm['rfm_score'] = rfm['recency_s'] * rfm['frequency_s'] * rfm['monetary_s']
```

```
[14]: rfm
```

```
[14]:
```

	recency	frequency	monetary	\
customer_unique_id				
0000366f3b9a7992bf8c76cfd3221e2	-112 days +19:55:50	1	141.90	

0000b849f77a49e4a4ce2b2a4ca5be3f	-115 days +20:10:50	1	27.19
0000f46a3911fa3c0805444483337064	-537 days +06:04:26	1	86.22
0000f6ccb0745a6a4b88665a16c9f078	-321 days +05:29:04	1	43.62
0004aac84e0df4da2b147fca70cf8255	-288 days +04:45:05	1	196.89
...
fffcf5a5ff07b0908bd4e2dbc735a684	-447 days +05:59:59	1	2067.42
fffea47cd6d3cc0a88bd621562a9d061	-262 days +05:07:19	1	84.58
ffff371b4d645b6ecea244b27531430a	-568 days +00:48:39	1	112.46
ffff5962728ec6157033ef9805bacc48	-119 days +00:17:04	1	133.69
ffffd2657e2aad2907e67c3e9daecbeb	-484 days +05:18:08	1	71.56

	recency_s	frequency_s	monetary_s	\
customer_unique_id				
0000366f3b9a7992bf8c76cfd3221e2	4	2	3	
0000b849f77a49e4a4ce2b2a4ca5be3f	4	2	1	
0000f46a3911fa3c0805444483337064	1	2	2	
0000f6ccb0745a6a4b88665a16c9f078	2	2	1	
0004aac84e0df4da2b147fca70cf8255	2	2	4	
...	
fffcf5a5ff07b0908bd4e2dbc735a684	1	2	5	
fffea47cd6d3cc0a88bd621562a9d061	2	2	2	
ffff371b4d645b6ecea244b27531430a	1	2	3	
ffff5962728ec6157033ef9805bacc48	4	2	3	
ffffd2657e2aad2907e67c3e9daecbeb	1	2	2	

	rfm_score
customer_unique_id	
0000366f3b9a7992bf8c76cfd3221e2	24
0000b849f77a49e4a4ce2b2a4ca5be3f	8
0000f46a3911fa3c0805444483337064	4
0000f6ccb0745a6a4b88665a16c9f078	4
0004aac84e0df4da2b147fca70cf8255	16
...	...
fffcf5a5ff07b0908bd4e2dbc735a684	10
fffea47cd6d3cc0a88bd621562a9d061	8
ffff371b4d645b6ecea244b27531430a	6
ffff5962728ec6157033ef9805bacc48	24
ffffd2657e2aad2907e67c3e9daecbeb	4

[93356 rows x 7 columns]

[]:

```
[15]: #rfm segment
start_time = timeit.default_timer()
rfm['segment'] = rfm['rfm_score'].apply(lambda x:
    ↳my_segment(rfm['rfm_score'],x,6))
```

```
elapsed = timeit.default_timer() - start_time

print(elapsed)
```

14.382687641998928

```
[16]: rfm['segment'].value_counts()
```

```
[16]: 1    76691
      2    11947
      3     3596
      4      510
      6      376
      5      236
      Name: segment, dtype: int64
```

```
[17]: fig2 =px.scatter_3d(rfm,
                        x='recency_s',
                        y='frequency_s',
                        z='monetary_s',
                        color='segment',
                        )
fig2.update_traces(marker=dict(size=3))
fig2.show()
```

Definition of the segments

Le *rfm_score* a été défini comme le produit des variables scores de chaque client. Ainsi l'aggrégation étant faite par multiplication, il y a une perte de l'information "spatiale" de chaque facteur. C'est pourquoi l'appréciation issue de ce score ne peut se faire que sur une échelle unique globale.

Pour illustrer cela considérons A et B deux clients. A et B n'ont passé en commandes que des produits haut de gamme. Ils ont donc un score monetary de 5. Mais si A l'a fait plusieurs fois - avec une satisfaction confirmée dans ses review_score, au moment de l'ouverture du magasin en ligne; B a commandé moins de fois, mais très récemment: ce dernier a donc pu être exposé à nos nouvelles collections ainsi qu'à notre dernière campagne marketing particulièrement bien performante considérant l'indicateur des liens de redirection facebook et youtube. Ils ont respectivement A(freq:4, recency:2), et B(freq:2, recency:4). Leur *rfm_score* égal de 40 exprime un **niveau global du potentiel de chiffre d'affaires**.

Segment 1: **l'étincelle**: *rfm_score* in [1,16]

Segment 2: **la buchette [d'allumette]**: *rfm_score* in [17,32]

Segment 3: **le briquet**: *rfm_score* in [33,48]

Segment 4: **le flambeau**: *rfm_score* in [49,64]

Segment 5: **le feu de camp**: *rfm_score* in [65,80]

Segment 6: **le dragon**: *rfm_score* in [81, 100]

0.0.2 Adding new variables related to the customers

Review score

```
[18]: part40 = reviews[['order_id', 'review_score']].groupby(['order_id']).mean()
order_review = orders[['order_id', 'customer_id']].merge(part40, on='order_id')
part41 = order_review[['customer_id', 'review_score']].
    ↳merge(customers[['customer_id', 'customer_unique_id']], on='customer_id')
part4 = part41[['customer_unique_id', 'review_score']].
    ↳groupby(['customer_unique_id']).mean()
```

Number of purchased products

```
[19]: part50 = items[['order_id', 'order_item_id']].groupby(['order_id']).sum()
order_items = orders[['order_id', 'customer_id']].merge(part50, on='order_id')
part51 = order_items[['customer_id', 'order_item_id']].
    ↳merge(customers[['customer_id', 'customer_unique_id']], on='customer_id')
part5 = part51[['customer_unique_id', 'order_item_id']].
    ↳groupby(['customer_unique_id']).sum()
```

```
[20]: rfm = (rfm.merge(part4, on='customer_unique_id')).merge(part5, on =
    ↳'customer_unique_id')
```

```
[21]: rfm = rfm.rename(columns={'order_item_id': 'num_item', 'review_score':
    ↳'review_s'})
```

```
[22]: score = rfm[['recency_s', 'monetary_s', 'review_s', 'num_item']]
```

```
[25]: rfm['recency_'] = rfm['recency'].apply(lambda x: to_seconds(x))
```

0.0.3 PCA run over the score dataframe

```
[26]: X = rfm[['recency_', 'monetary', 'review_s', 'num_item']]
X_scaled = preprocessing.StandardScaler().fit_transform(X)
s_scaled = pd.DataFrame(X_scaled, columns=X.columns, index=X.index)

#Display
s_scaled.describe()
```

```
[26]:
```

	recency_	monetary	review_s	num_item
count	8.856800e+04	8.856800e+04	8.856800e+04	8.856800e+04
mean	2.058290e-17	1.599656e-16	1.191988e-15	1.156105e-15
std	1.000006e+00	1.000006e+00	1.000006e+00	1.000006e+00
min	-3.115415e+00	-7.297696e-01	-2.427112e+00	-1.727628e-01
25%	-7.111577e-01	-4.588921e-01	-1.072805e-01	-1.727628e-01
50%	1.235022e-01	-2.520014e-01	6.659967e-01	-1.727628e-01


```

75%    8.097209e-01  8.254519e-02  6.659967e-01 -1.727628e-01
max    1.555779e+00  3.363083e+01  6.659967e-01  9.634392e+01

```

```

[27]: pca = decomposition.PCA(n_components=4)
      n_comp=4
      X_projected = pca.fit_transform(s_scaled)
      score_pc = pd.DataFrame(X_projected, index=score.index, columns=["F"+str(i+1)
      ↪for i in range(n_comp)])

      #Display
      score_pc

```

```

[27]:

```

	F1	F2	F3	F4
customer_unique_id				
0000366f3b9a7992bf8c76cfd3221e2	-0.405310	0.962735	0.270200	0.038669
0000b849f77a49e4a4ce2b2a4ca5be3f	-0.452033	0.670076	-0.588078	0.250455
0000f46a3911fa3c0805444483337064	-0.036028	-2.159198	-0.312470	-0.029606
0000f6ccb0745a6a4b88665a16c9f078	-0.424089	-0.610754	-0.146662	0.217431
0004aac84e0df4da2b147fca70cf8255	-0.263005	-0.113414	0.716821	-0.113463
...
ffffcf5a5ff07b0908bd4e2dbc735a684	5.699430	-0.210566	4.306526	-5.153173
fffea47cd6d3cc0a88bd621562a9d061	-0.300039	-0.224101	-0.195463	0.088717
ffff371b4d645b6ecea244b27531430a	-0.533729	-1.899279	1.134881	0.164665
ffff5962728ec6157033ef9805bacc48	-0.429795	0.910205	0.271965	0.064145
ffffd2657e2aad2907e67c3e9daecbeb	-0.644042	-1.393117	0.897163	0.281629

```

[88568 rows x 4 columns]

```


```

[28]: ##Definition de la fonction d'affichage des valeurs propres
def display_screes_plot(pca):
    scree = pca.explained_variance_ratio_*100
    plt.bar(np.arange(len(scree))+1, scree);
    plt.plot(np.arange(len(scree))+1, scree.cumsum(),c="red",marker='o');
    plt.xlabel("rank of the axis of inertia");
    plt.ylabel("percentage of inertia");
    plt.title("Decay of eigen-values");

    plt.savefig('./EigenValuesDecay.png', bbox_inches = 'tight');
    plt.show(block=False);

    ##Tracé de l'éboulis
    display_screes_plot(pca)

```



```
[29]: ##Definition de la fonction d'affichage des valeurs propres
dic_graph={}
def display_circles(pcs, n_comp, pca, axis_ranks, labels=None,
    ↪label_rotation=0, lims=None):
    for d1, d2 in axis_ranks: # On affiche les 3 premiers plans factoriels,
    ↪donc les 6 premières composantes
        if d2 < n_comp:

            # initialisation de la figure
            fig, ax = plt.subplots(figsize=(7,6))

            # détermination des limites du graphique
            if lims is not None :
                xmin, xmax, ymin, ymax = lims
            elif pcs.shape[1] < 30 :
                xmin, xmax, ymin, ymax = -1, 1, -1, 1
            else :
                xmin, xmax, ymin, ymax = min(pcs[d1,:]), max(pcs[d1,:]),
    ↪min(pcs[d2,:]), max(pcs[d2,:])

            # affichage des flèches
            # s'il y a plus de 30 flèches, on n'affiche pas le triangle à leur
    ↪extrémité
```

```

if pcs.shape[1] < 30 :
    plt.quiver(np.zeros(pcs.shape[1]), np.zeros(pcs.shape[1]),
               pcs[d1,:], pcs[d2,:],
               angles='xy', scale_units='xy', scale=1, color="grey")
    # (voir la doc : https://matplotlib.org/api/\_as\_gen/matplotlib.pyplot.quiver.html)
else:
    lines = [[[0,0],[x,y]] for x,y in pcs[[d1,d2]].T]
    ax.add_collection(LineCollection(lines, axes=ax, alpha=.1,
    ↪color='black'))

    # affichage des noms des variables
    if labels is not None:
        for i,(x, y) in enumerate(pcs[[d1,d2]].T):
            if x >= xmin and x <= xmax and y >= ymin and y <= ymax :
                plt.text(x, y, labels[i], fontsize='14', ha='center',
    ↪va='center', rotation=label_rotation, color="blue", alpha=0.5)

    # affichage du cercle
    circle = plt.Circle((0,0), 1, facecolor='none', edgecolor='b')
    plt.gca().add_artist(circle)

    # définition des limites du graphique
    plt.xlim(xmin, xmax)
    plt.ylim(ymin, ymax)

    # affichage des lignes horizontales et verticales
    plt.plot([-1, 1], [0, 0], color='grey', ls='--')
    plt.plot([0, 0], [-1, 1], color='grey', ls='--')

    # nom des axes, avec le pourcentage d'inertie expliqué
    plt.xlabel('F{} ({}%)'.format(d1+1, round(100*pca.
    ↪explained_variance_ratio_[d1],1)))
    plt.ylabel('F{} ({}%)'.format(d2+1, round(100*pca.
    ↪explained_variance_ratio_[d2],1)))

    plt.title("Cercle of correlations (F{} et F{}).format(d1+1, d2+1))

    # Enregistrement du tracé du cercle de corr de ce plan factoriel
    plt.savefig('./CercleCorr'+str(d1)+str(d2)+'.png',
    ↪bbox_inches='tight')


plt.show(block=False)

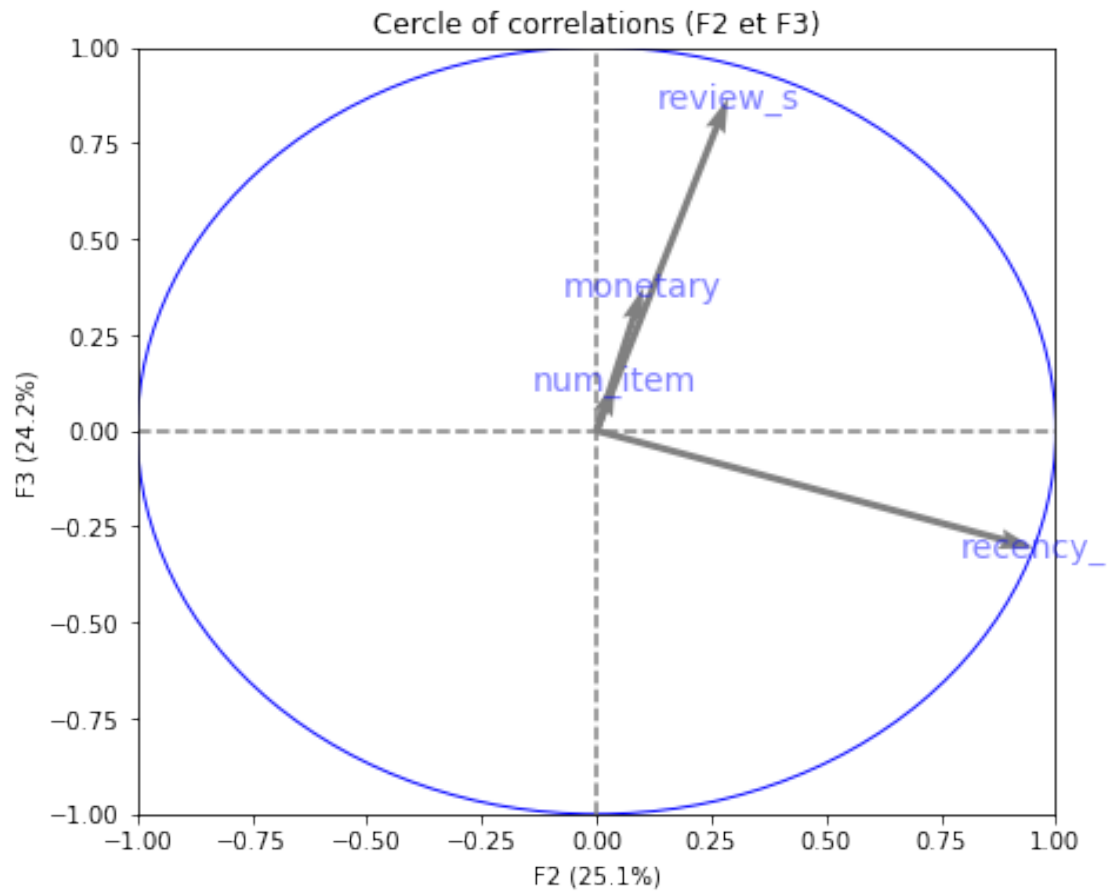
```

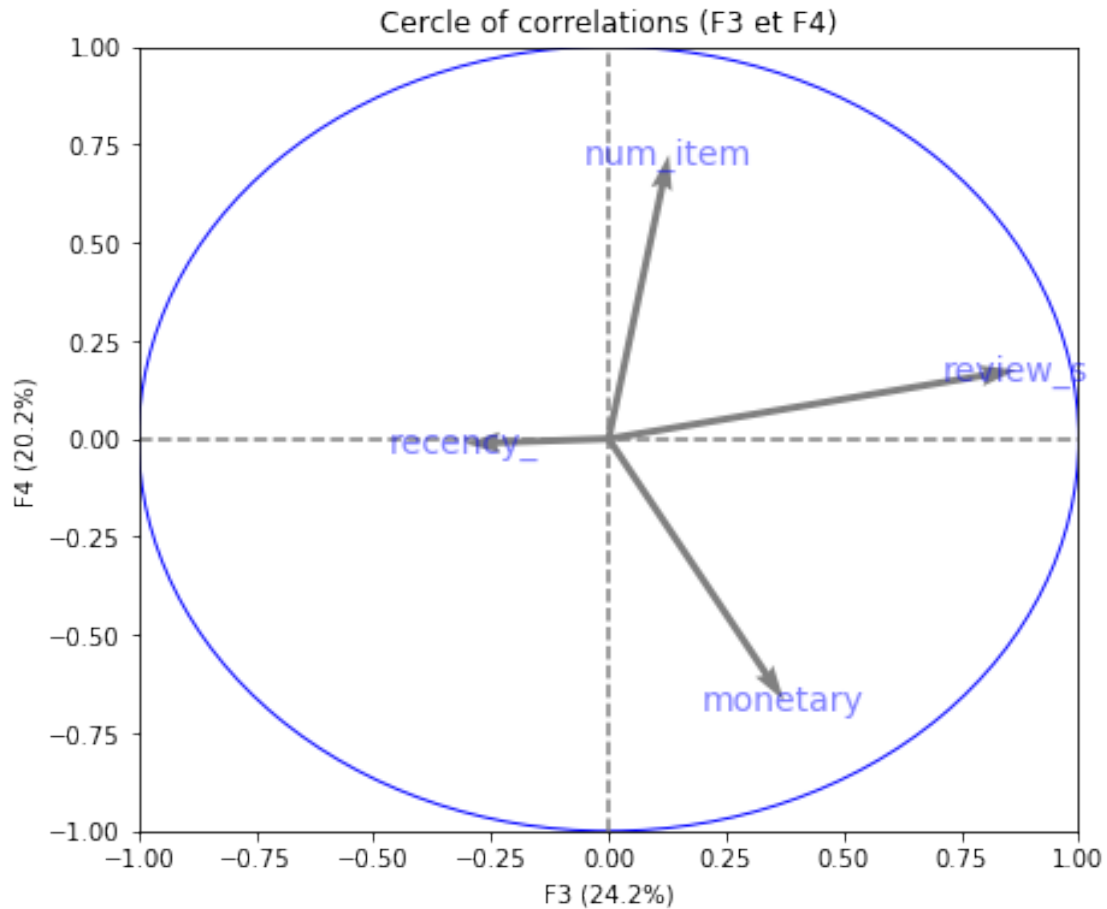
```
##Tracé du cercle des corrélations
```

```
pcs = pca.components_
```

```
display_circles(pcs, 4, pca, [(0,1),(1,2),(2,3)], labels=s_scaled.columns)
```







```
[30]: ##Definition de la fonction d'affichage des valeurs propres
def display_factorial_planes(X_projected, n_comp, pca, axis_ranks, labels=None,
    ↪ alpha=0.5, illustrative_var=None):
    for d1,d2 in axis_ranks:
        if d2 < n_comp:

            # initialisation de la figure
            fig = plt.figure(figsize=(7,6))

            # affichage des points
            if illustrative_var is None:
                plt.scatter(X_projected[:, d1], X_projected[:, d2], alpha=alpha)
            else:
                illustrative_var = np.array(illustrative_var)
                for value in np.unique(illustrative_var):
                    selected = np.where(illustrative_var == value)
                    plt.scatter(X_projected[selected, d1],
    ↪ X_projected[selected, d2], alpha=alpha, label=value)
```

```

plt.legend()

# affichage des labels des points
if labels is not None:
    for i,(x,y) in enumerate(X_projected[:, [d1,d2]]):
        plt.text(x, y, labels[i],
                 fontsize='14', ha='center', va='center')

# détermination des limites du graphique
boundary = np.max(np.abs(X_projected[:, [d1,d2]])) * 1.1
plt.xlim([-boundary,boundary])
# plt.ylim([-boundary,boundary])
plt.ylim([-0.1,boundary])

# affichage des lignes horizontales et verticales
plt.plot([-100, 100], [0, 0], color='grey', ls='--')
plt.plot([0, 0], [-100, 100], color='grey', ls='--')

# nom des axes, avec le pourcentage d'inertie expliqué
plt.xlabel('F{} ({}%)'.format(d1+1, round(100*pca.
→explained_variance_ratio_[d1],1)))
plt.ylabel('F{} ({}%)'.format(d2+1, round(100*pca.
→explained_variance_ratio_[d2],1)))

plt.title("Projection of individuals (on F{} et F{})."format(d1+1,
→d2+1))

#Enregistrement du tracé
plt.savefig('./ProjectionOn'+str(d1)+str(d2), bbox_inches='tight')

plt.show(block=False)

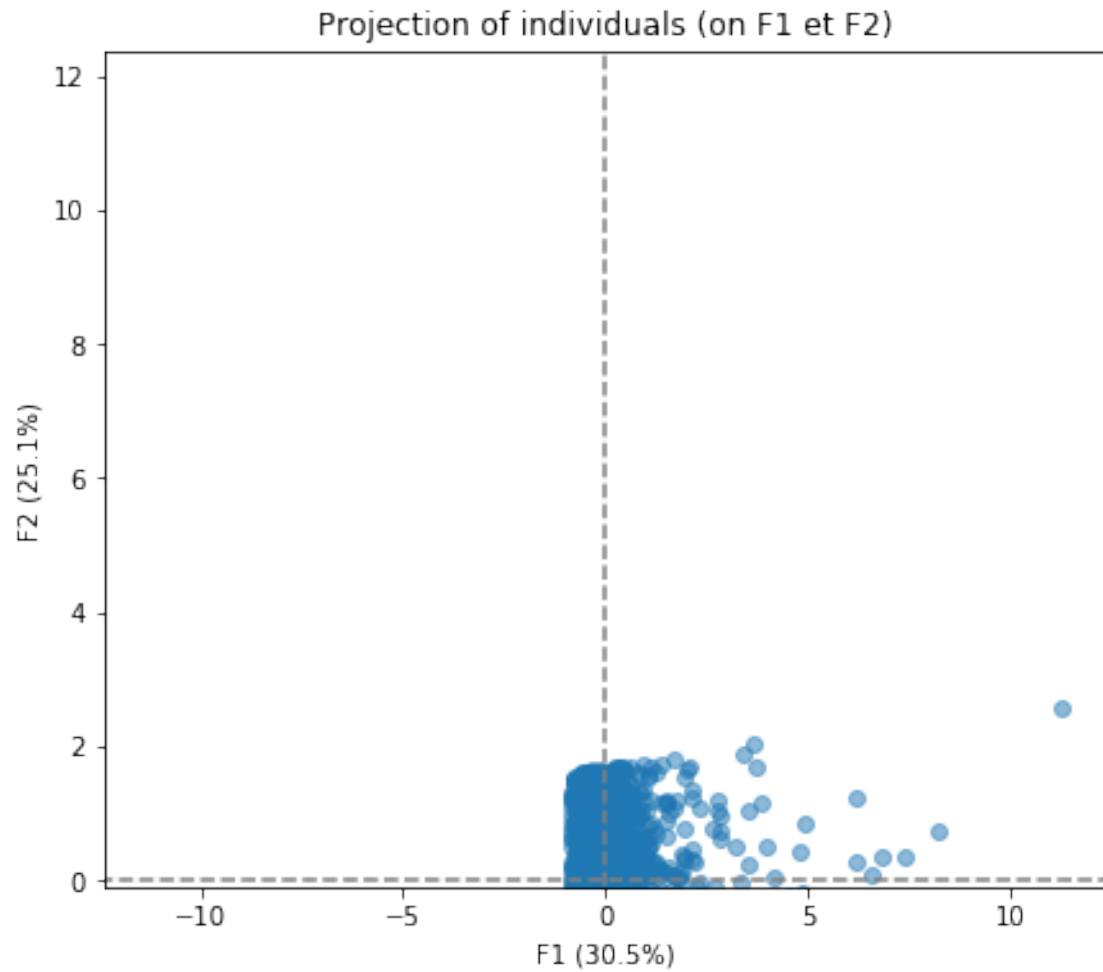
##Tracé du cercle des corrélations
pcs = pca.components_
sampled = np.random.choice(range(X_projected.shape[0]), 2000, replace=False) #!
→rend une liste d'indices

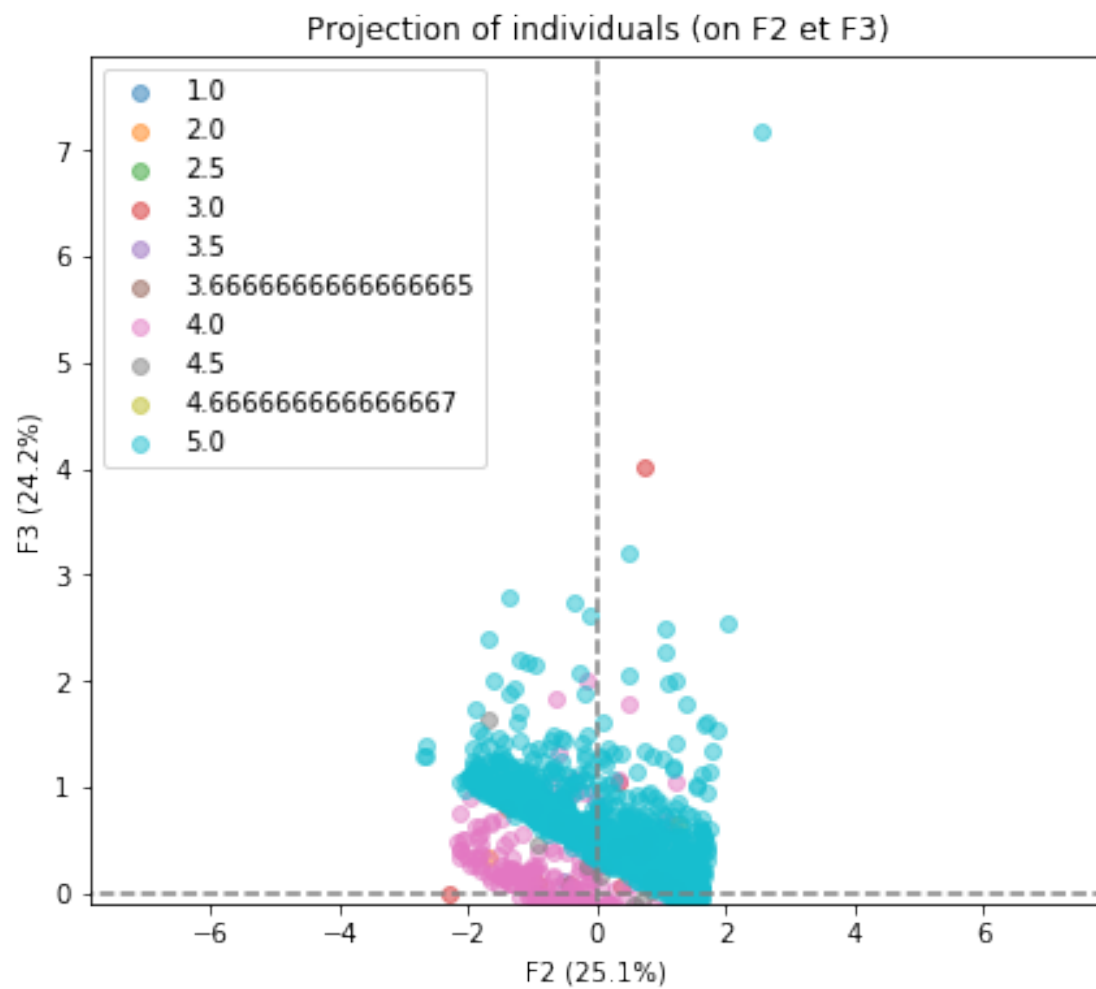
display_factorial_planes(X_projected[sampled,:], n_comp, pca, [(0,1)])
display_factorial_planes(X_projected[sampled,:], n_comp, pca, [(1,2)],
→illustrative_var=X['review_s'].iloc[sampled])

#display_factorial_planes(X_projected[sampled,:], n_comp, pca, [(0,1)],
→illustrative_var=X['recency_'].iloc[sampled])
#display_factorial_planes(X_projected[sampled,:], n_comp, pca, [(0,1)],
→illustrative_var=X['monetary'].iloc[sampled])

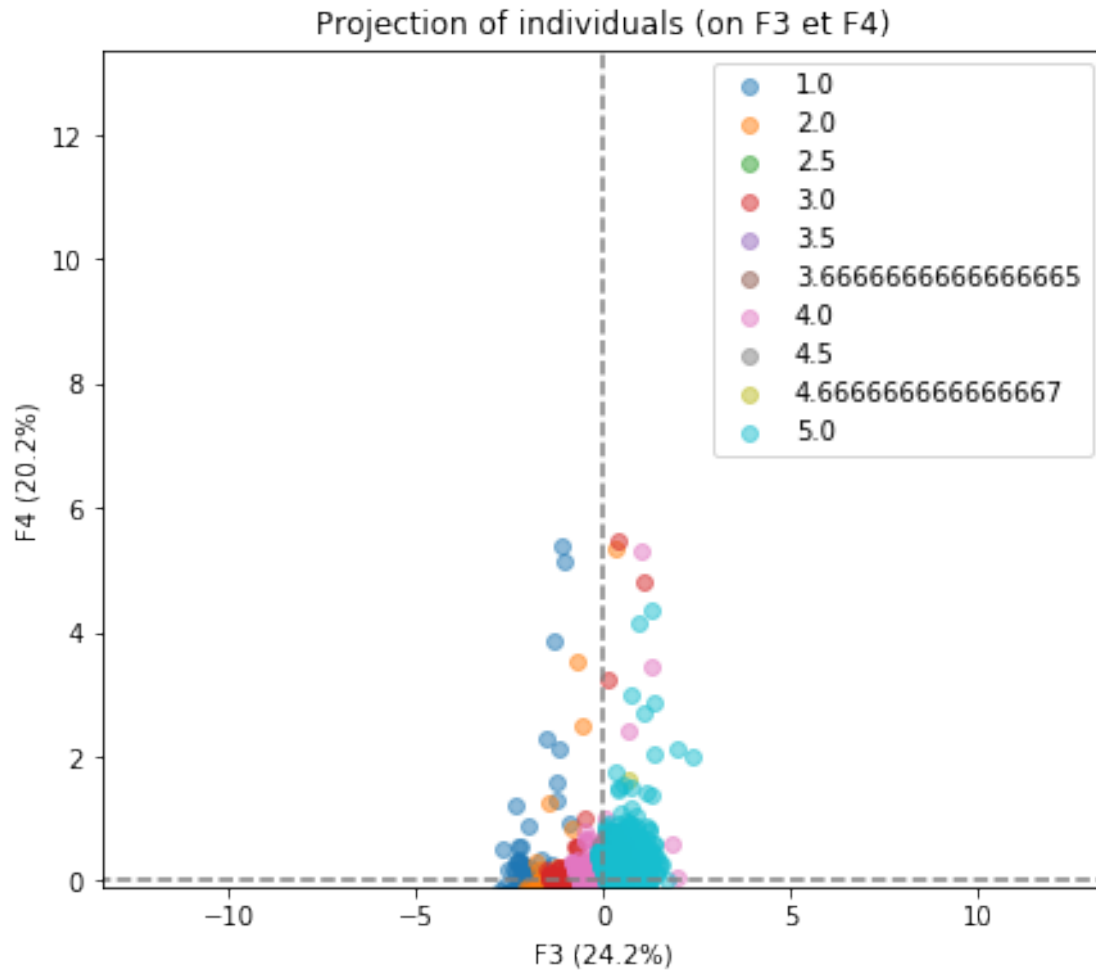
```

```
#display_factorial_planes(X_projected[sampled,:], n_comp, pca, [(0,1)],  
→ illustrative_var=X['review_s'].iloc[sampled])  
#display_factorial_planes(X_projected[sampled,:], n_comp, pca, [(0,1)],  
→ illustrative_var=X['num_item'].iloc[sampled])
```





```
[31]: display_factorial_planes(X_projected[sampled,:], n_comp, pca, [(2,3)],  
    ↪ illustrative_var=X['review_s'].iloc[sampled])
```



0.0.4 Kmeans based on silhouette score

```
[32]: #Extraction of the 2 first latent features
```

```
latent_features = score_pc[['F2', 'F3']]
#X = latent_features[sampled,:]
```

```
X = latent_features
```

```
[ ]:
```

```
[33]: range_n_clusters = [2, 3, 4, 5, 6]
```

```
for n_clusters in range_n_clusters:
    # Create a subplot with 1 row and 2 columns
    fig, (ax1, ax2) = plt.subplots(1, 2)
```

```

fig.set_size_inches(18, 7)

# The 1st subplot is the silhouette plot
# The silhouette coefficient can range from -1, 1 but in this example all
# lie within [-0.1, 1]
ax1.set_xlim([-0.1, 1])
# The (n_clusters+1)*10 is for inserting blank space between silhouette
# plots of individual clusters, to demarcate them clearly.
ax1.set_ylim([0, len(X) + (n_clusters + 1) * 10])

# Initialize the clusterer with n_clusters value and a random generator
# seed of 10 for reproducibility.
clusterer = KMeans(n_clusters=n_clusters, random_state=10)
cluster_labels = clusterer.fit_predict(X)

# The silhouette_score gives the average value for all the samples.
# This gives a perspective into the density and separation of the formed
# clusters
start_time = timeit.default_timer()
silhouette_avg = silhouette_score(X, cluster_labels, sample_size=10000)
→ #Used argument sample_size on Dec 28, 2020
elapsed = timeit.default_timer() - start_time
print("For n_clusters =", n_clusters,
      "The average silhouette_score is :", silhouette_avg, "and its
→ computation time is: ", elapsed) # time computation added on Dec 28, 2020

# Compute the silhouette scores for each sample
sample_silhouette_values = silhouette_samples(X, cluster_labels)

y_lower = 10
for i in range(n_clusters):
    # Aggregate the silhouette scores for samples belonging to
    # cluster i, and sort them
    ith_cluster_silhouette_values = \
        sample_silhouette_values[cluster_labels == i]

    ith_cluster_silhouette_values.sort()

    size_cluster_i = ith_cluster_silhouette_values.shape[0]
    y_upper = y_lower + size_cluster_i

    color = cm.nipy_spectral(float(i) / n_clusters)
    ax1.fill_betweenx(np.arange(y_lower, y_upper),
                      0, ith_cluster_silhouette_values,
                      facecolor=color, edgecolor=color, alpha=0.7)

    # Label the silhouette plots with their cluster numbers at the middle

```

```

ax1.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))

# Compute the new y_lower for next plot
y_lower = y_upper + 10 # 10 for the 0 samples

ax1.set_title("The silhouette plot for the various clusters.")
ax1.set_xlabel("The silhouette coefficient values")
ax1.set_ylabel("Cluster label")

# The vertical line for average silhouette score of all the values
ax1.axvline(x=silhouette_avg, color="red", linestyle="--")

ax1.set_yticks([]) # Clear the yaxis labels / ticks
ax1.set_xticks([-0.1, 0, 0.2, 0.4, 0.6, 0.8, 1])

# 2nd Plot showing the actual clusters formed
colors = cm.nipy_spectral(cluster_labels.astype(float) / n_clusters)
ax2.scatter(X.iloc[:, 0], X.iloc[:, 1], marker='.', s=30, lw=0, alpha=0.7,
            c=colors, edgecolor='k')

# Labeling the clusters
centers = clusterer.cluster_centers_
# Draw white circles at cluster centers
ax2.scatter(centers[:, 0], centers[:, 1], marker='o',
            c="white", alpha=1, s=200, edgecolor='k')

for i, c in enumerate(centers):
    ax2.scatter(c[0], c[1], marker='$%d$' % i, alpha=1,
                s=50, edgecolor='k')

ax2.set_title("The visualization of the clustered data.")
ax2.set_xlabel("Feature space for the 1st feature")
ax2.set_ylabel("Feature space for the 2nd feature")

plt.suptitle(("Silhouette analysis for KMeans clustering on sample data "
            "with n_clusters = %d" % n_clusters),
            fontsize=14, fontweight='bold')

plt.savefig('./HistoPayValues_%d.png' % n_clusters, bbox_inches='tight');

plt.show()

```

For n_clusters = 2 The average silhouette_score is : 0.37268564662444226 and its computation time is: 0.8707660030049738

For n_clusters = 3 The average silhouette_score is : 0.4588228312073049 and its computation time is: 0.9076170300031663



For n_clusters = 4 The average silhouette_score is : 0.3982032764823717 and its

computation time is: 0.8008753609974519



For $n_clusters = 5$ The average silhouette_score is : 0.4054340042179942 and its computation time is: 0.795828957001504

For $n_clusters = 6$ The average silhouette_score is : 0.40478601244076334 and its computation time is: 0.8044797370021115



Silhouette analysis for KMeans clustering on sample data with $n_clusters = 2$





Silhouette analysis for KMeans clustering on sample data with $n_clusters = 3$




Silhouette analysis for KMeans clustering on sample data with $n_clusters = 4$



Silhouette analysis for KMeans clustering on sample data with $n_clusters = 5$





```
[34]: km_all = rfm[['recency_', 'monetary', 'review_s', 'num_item']]
      #km_all = latent_features
      km_all['segment'] = cluster_labels

      km_seg_counts = km_all.groupby(['segment']).count()
      km_seg_mean = km_all.groupby(['segment']).mean()
      km_seg_var = km_all.groupby(['segment']).var()
```

/home/erbadi/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3:
SettingWithCopyWarning:

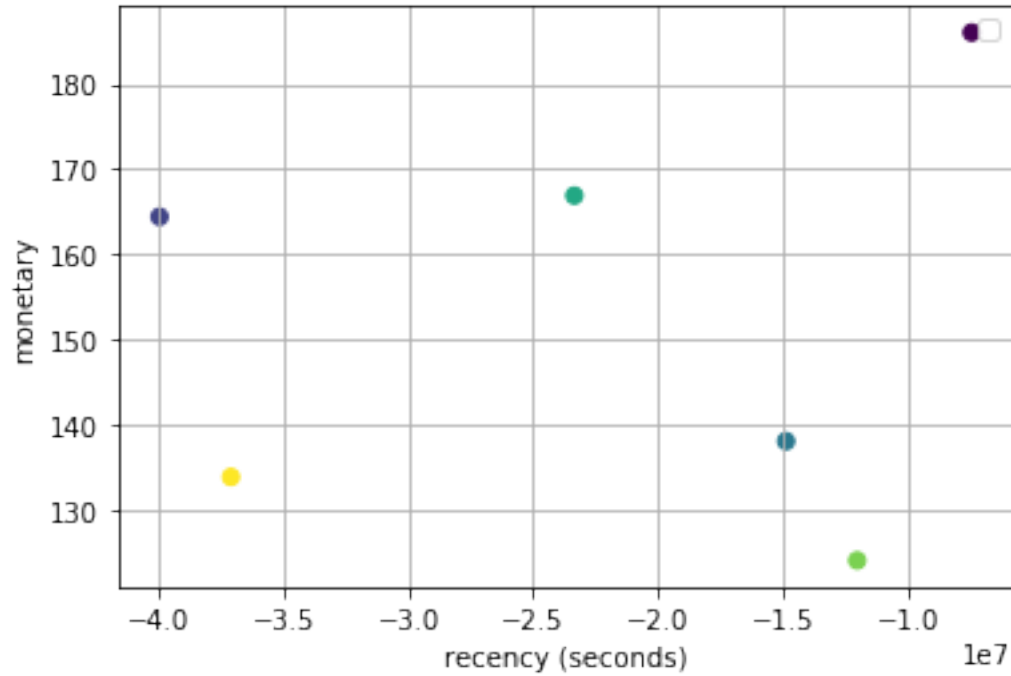
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
[35]: plt.scatter(km_seg_mean['recency_'], km_seg_mean['monetary'], c=km_seg_mean.
      ↪index);
      plt.xlabel('recency (seconds)')
      plt.ylabel('monetary')
      plt.legend()
      plt.grid()

      plt.savefig('./Centroids', bbox_inches='tight')
```

No handles with labels found to put in legend.



[36]: km_seg_mean

[36]:

	recency_	monetary	review_s	num_item
segment				
0	-7.476280e+06	186.035972	4.919927	1.473606
1	-3.992703e+07	164.432451	4.703270	1.347842
2	-1.489350e+07	138.111701	1.193843	1.522264
3	-2.335253e+07	166.913704	4.775549	1.389434
4	-1.205166e+07	124.132501	3.529963	1.331907
5	-3.707289e+07	133.917814	2.048375	1.455801

Segment 0: le dragon déchainé

Segment 1: le phoenix endormi

Segment 2: la braise fraîche

Segment 3: le feu de camp interrompu

Segment 4: l'étincelle

Segment 5: l'allumette fumante


0.0.5 DBScan segmentation

Determination of the optimal epsilon


```
[37]: %matplotlib inline
from sklearn.neighbors import NearestNeighbors
import matplotlib.pyplot as plt
import numpy as np

neigh = NearestNeighbors(n_neighbors=5)
nbrs = neigh.fit(score_pc)
distances, indices = nbrs.kneighbors(score_pc)
distances = np.sort(distances, axis=0)
distances = distances[:,1]
plt.plot(distances);
plt.semilogy()
plt.xlabel('individuals')
plt.ylabel('distance')
plt.title('Distance to the nearest neighbor')

plt.savefig('./NearestNeighbor.png', bbox_inches='tight')
```



Une bonne valeur de epsilon serait donc **0.1**


Segmentation

```
[38]: from sklearn.cluster import DBSCAN


start_time = timeit.default_timer()
y_pred = DBSCAN(eps = 0.1, min_samples=5).fit_predict(score_pc)
elapsed = timeit.default_timer() - start_time
print(elapsed)

plt.scatter(score_pc['F2'],score_pc['F3'],c = y_pred);
```

1.7405096310030785



```
[39]: plt.scatter(score_pc['F1'],score_pc['F2'],c = y_pred);
```



Interpretation

```
[40]: #Extracting the means and sqrt of each segment
#db_all = score_pc[['F2', 'F3']]
db_all = rfm[['recency_', 'monetary_', 'review_s', 'num_item']]
db_all['segment']=y_pred

seg_counts = db_all.groupby(['segment']).count()
seg_mean = db_all.groupby(['segment']).mean()
seg_sqrt = db_all.groupby(['segment']).var()

thresh = 1000
populated_seg = seg_mean.where(seg_counts['recency_']>=thresh)
#db_all.groupby(['segment']).mean()
plt.scatter(populated_seg['recency_'],populated_seg['monetary_'],c=seg_counts['recency_']);
plt.xlabel('recency_')
plt.ylabel('monetary_')
plt.title('Centers of clusters counting more than %d individuals' % thresh )

plt.savefig('./Clusters_%d.png' % thresh, bbox_inches='tight');
print('Le nombre de clients couverts par les segments ayant un effectif de plus de: ',thresh, ' est: ',seg_counts[seg_counts['recency_']>=thresh]['recency_'].sum())
```


Le nombre de clients couverts par les segments ayant un effectif de plus de:

1000 est: 83911

/home/erbadi/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:4:
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy



0.0.6 Segment stability

New version

```
[41]: # Function which updates the period considered and then computes the different_
      ↪ aggregated variables used for the subsequent segmentation
def extended_basis(order, b, timestep, init='recency'):
    #Updating the period covered
    part0 = order[order['order_purchase_timestamp'] <= b[init].max() + datetime.
    ↪timedelta(weeks=+timestep)]
```

```

#Processing related to "orders"
order_custom = part0.merge(customers[['customer_id',
→ 'customer_unique_id']], on='customer_id')
part1 = order_custom[['order_purchase_timestamp', 'customer_unique_id']].
→groupby(['customer_unique_id']).max()

#Processing related to "customers"
# part2 = customers[['customer_id', 'customer_unique_id']].
→groupby(['customer_unique_id']).count()

#Processing related to "payments"
part30 = payments[['payment_value', 'order_id']].groupby(['order_id']).sum()
order_payment = part0.join(part30, on=['order_id'])
part31 = order_payment[['payment_value', 'customer_id']].
→merge(customers[['customer_unique_id', 'customer_id']], on='customer_id')
part3 = part31[['payment_value', 'customer_unique_id']].
→groupby(['customer_unique_id']).sum()

#Processing related to reviews
part40 = reviews[['order_id', 'review_score']].groupby(['order_id']).mean()
order_review = part0[['order_id', 'customer_id']].merge(part40,
→on='order_id')
part41 = order_review[['customer_id', 'review_score']].
→merge(customers[['customer_id', 'customer_unique_id']], on='customer_id')
part4 = part41[['customer_unique_id', 'review_score']].
→groupby(['customer_unique_id']).mean()

#Processing related to items [purchased]
part50 = items[['order_id', 'order_item_id']].groupby(['order_id']).sum()
order_items = part0[['order_id', 'customer_id']].merge(part50, on='order_id')
part51 = order_items[['customer_id', 'order_item_id']].
→merge(customers[['customer_id', 'customer_unique_id']], on='customer_id')
part5 = part51[['customer_unique_id', 'order_item_id']].
→groupby(['customer_unique_id']).sum()

rfm = ((part1.merge(part3, on='customer_unique_id')).merge(part4, on =
→ 'customer_unique_id')).merge(part5, on = 'customer_unique_id')
rfm = rfm.rename(columns={'order_purchase_timestamp':
→ 'recency', 'payment_value': 'monetary', 'order_item_id': 'num_item',
→ 'review_score': 'review_s'})
rfm['recency'] = rfm['recency'] - order['order_purchase_timestamp'].max()
res = rfm

return res

```

```

# Function used to compute a scaling of the data, in order for them to match
↳with the '.fit(X)' arguments requirements
def scaled_basis(b, fitted_scaler):
    X = b[['monetary', 'review_s', 'num_item']]
    X['recency_'] = b['recency'].apply(lambda x: to_seconds(x))
    X_scaled = fitted_scaler.transform(X)
    b_scaled = pd.DataFrame(X_scaled, columns=X.columns, index=X.index)
    return b_scaled

```

```

[42]: def stability_check(data, b0, n_clusters, timestep, num_timestep):

    #Building and fitting the scaler of the bases
    X = b0[['monetary', 'review_s', 'num_item']]
    X['recency_'] = b0['recency'].apply(lambda x: to_seconds(x))
    fitted_scaler = preprocessing.StandardScaler().fit(X)

    #Building and fitting the clusterer
    clusterer = KMeans(n_clusters=n_clusters, random_state=10)
    c0 = clusterer.fit(scaled_basis(b0, fitted_scaler))

    bases = {'b0':b0}
    clusterers = {'c0':c0}
    current_clusterings = {}
    updated_clusterings = {}
    rand_scores = {}
    for ind in range(num_timestep):
        #Future bases and forecast clusterings (with current clusterer)
        bases['b'+str(ind+1)] =
↳extended_basis(data, bases['b'+str(ind)], timestep)
        current_clusterings['c0.'+str(ind+1)] = c0.
↳predict(scaled_basis(bases['b'+str(ind+1)], fitted_scaler))

        #Updated clusterings (-ers)
        clusterer = KMeans(n_clusters=n_clusters, random_state=10)
        clusterers['c'+str(ind+1)] = clusterer.
↳fit(scaled_basis(bases['b'+str(ind+1)], fitted_scaler))
        updated_clusterings['c'+str(ind+1)+'.'+str(ind+1)] =
↳clusterers['c'+str(ind+1)].predict(scaled_basis(bases['b'+str(ind+1)],
↳fitted_scaler))

        #Rand score computation
        rand_scores[ind+1] = adjusted_rand_score(current_clusterings['c0.
↳'+str(ind+1)],
        ↳
↳updated_clusterings['c'+str(ind+1)+'.'+str(ind+1)])

```

```

#Plotting the rand score evolution
lists = sorted(rand_scores.items()) # sorted by key, return a list of tuples
x, y = zip(*lists) # unpack a list of pairs into two tuples
plt.plot(x, y)
plt.xlabel('Step')
plt.ylabel('Rand score')
plt.title('Randscore decay - Forecast vs future segmentation')

plt.savefig('./Ebouli_Randscore.png', bbox_inches='tight')
plt.show()

#return the rand score
return bases

```

Actual computation

```

[43]: #Adjustment of the dataset orders to the needs of the computation -> from
      ↪ datetime to timedelta by recentering
orders['order_purchase_timestamp'] = orders['order_purchase_timestamp'] -
      ↪ orders['order_purchase_timestamp'].max()

```

```

[44]: #Parameters for the computation of the segments stability assessment
data = orders
init_time = orders['order_purchase_timestamp'].min()
init_b = orders[orders['order_purchase_timestamp']== init_time]

b0 = extended_basis(orders, init_b, 52, init='order_purchase_timestamp')
#b0 = data[data['recency'] <= data['recency'].min() + datetime.
      ↪ timedelta(weeks=+52)]
n_clusters = 6
timestep = 4 #time step in weeks unit
num_timestep = 14

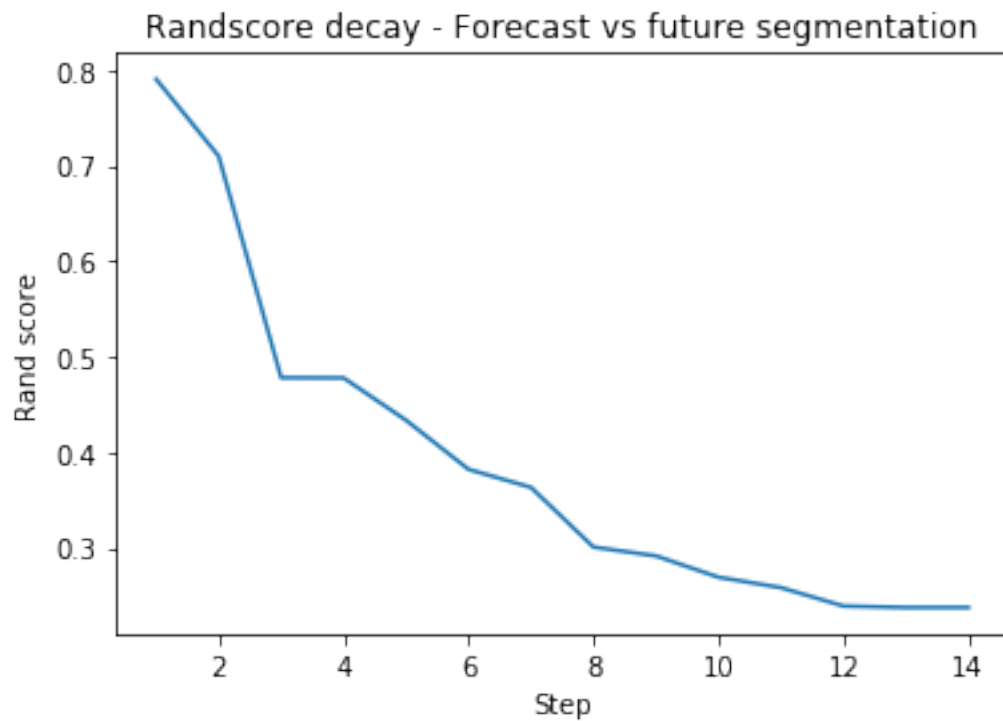
```

```

[45]: start_time = timeit.default_timer()
res = stability_check(data, b0, n_clusters, timestep, num_timestep)
elapsed = timeit.default_timer() - start_time

print('The computation time is: ', elapsed)

```



The computation time is: 40.18205604499963

```
[46]: for ind in list(res.values()):  
      print(ind.shape)
```

```
(22337, 4)  
(25883, 4)  
(29548, 4)  
(36680, 4)  
(40760, 4)  
(46771, 4)  
(52748, 4)  
(58681, 4)  
(64428, 4)  
(70660, 4)  
(75284, 4)  
(79927, 4)  
(86689, 4)  
(88568, 4)  
(88568, 4)
```


1 SS

$$[\]:$$

[47] : b0

```
[47]: recency    monetary    review_s    \
```

customer_unique_id

```
0000f46a3911fa3c0805444483337064 -537 days +06:04:26      86.22      3.0
```

```
0005e1862207bf6ccc02e4228effd9a0 -543 days +08:31:35    150.12    4.0
```

0006fdc98a402fceb4eb0ee528f6a8d4	-408 days +18:22:33	29.00	3.0
----------------------------------	---------------------	-------	-----

000a5ad9c4601d2bbdd9ed765d5213b3	-384 days +22:44:38	91.28	4.0
----------------------------------	---------------------	-------	-----

```
000de6019bb59f34c099a907c151d855 -377 days +04:09:56 257.44 2.0
```

...

```
fff3a9369e4b7102fab406a334a678c3 -384 days +19:26:01 102.74 5.0
```

```
fff699c184bcc967d62fa2c6171765f7 -362 days +02:06:17 55.00 4.0
```

```
fffcf5a5ff07b0908bd4e2dbc735a684 -447 days +05:59:59 2067.42 5.0
```

```
ffff371b4d645b6ecea244b27531430a -568 days +00:48:39 112.46 5.0
```

```
ffffd2657e2aad2907e67c3e9daecbeb -484 days +05:18:08    71.56    5.0
```

```
num_item
```

customer_unique_id

```
0000f46a3911fa3c0805444483337064 1
```

0005e1862207bf6ccc02e4228effd9a0	1
----------------------------------	---

```
0006fdc98a402fceb4eb0ee528f6a8d4      1
```

000a5ad9c4601d2bbdd9ed765d5213b3	1
----------------------------------	---

```
000de6019bb59f34c099a907c151d855 3
```

...

```
fff3a9369e4b7102fab406a334a678c3      1
```

```
fff699c184bcc967d62fa2c6171765f7 1
```

```
ffffcf5a5ff07b0908bd4e2dbc735a684 3
```

```
ffff371b4d645b6ecea244b27531430a 1
```

```
ffffd2657e2aad2907e67c3e9daecbeb 1
```

```
[22337 rows x 4 columns]
```